

Probabilistic program induction of symbols

Or: More complex models!

- *The Language of Thought: computational cognitive science approaches to category learning*
- Who: Fausto Carcassi
- When: Sommer semester 2022

Where are we?

- Last week, we have seen two applications of the idea of pLoT to new conceptual domains: handwritten digits and kinship systems.
- This week, we'll look at a few more applications.
- It'll be a bit of a whirlwind!
- We will look at:
 - Learning numerals in an LoT
 - Learning abstract visual concepts in an LoT
 - Learning sequences in an LoT

Learning Numerals with an LoT

- Piantadosi et al (2012), *Bootstrapping in a language of thought: A formal model of numerical concept learning*.
- Children exhibit very regular patterns in the way they learn number systems.
- The first learn to recognize small sets of size 1, then 2, then 3, etc.
 - In Carey's formulation, early number-word meanings are represented using mental models of small sets. For instance two-knowers might have a mental model of "one" as $\{X\}$ and "two" as $\{X, X\}$. These representations rely on children's ability for enriched parallel individuation, a representational capacity that Le Corre and Carey (2007) argue can individuate objects, manipulate sets, and compare sets using one-to-one correspondence. Subset-knowers can, for instance, check if "two" applies to a set S by seeing if S can be put in one-to-one correspondence with their mental model of two, $\{X, X\}$
- Then at about 3;6 they learn the full recursive system (Cardinal Principal learners)
- This is a qualitative jump rather than continuous smooth progress.

Learning Numerals with an LoT

- Let's see if we can reproduce this qualitative conceptual jump with an LoT!
- We have at least three choices for how to set up the LoT. Each sentence in the LoT could be:
 - A function from a number word to a predicate of sets
 - A function from a set to a number word
 - A function that constructs a set from the objects in the situation
- Children can do all these three things, and there is no clear empirical evidence one way or the other.
- In the paper, they go the second way: from a set to a number word.

Learning Numerals with an LoT

- Rules in the LoT:

Functions mapping sets to truth values

(singleton? X)

Returns true iff the set X has exactly one element

(doubleton? X)

Returns true iff the set X has exactly two elements

(tripleton? X)

Returns true iff the set X has exactly three elements

Functions on sets

(set-difference X Y)

Returns the set that results from removing Y from X

(union X Y)

Returns the union of sets X and Y

(intersection X Y)

Returns the intersect of sets X and Y

(select X)

Returns a set containing a single element from X

Logical functions

(and P Q)

Returns TRUE if P and Q are both true

(or P Q)

Returns TRUE if either P or Q is true

(not P)

Returns TRUE iff P is false

(if P X Y)

Returns X iff P is true, Y otherwise

Functions on the counting routine

(next W)

Returns the word after W in the counting routine

(prev W)

Returns the word before W in the counting routine

(equal-word? W V)

Returns TRUE if W and V are the same word

Recursion

(L S)

Returns the result of evaluating the entire current lambda expression on set S

Learning Numerals with an LoT

- Recursion is the only one that's a bit complicated. What do you think the following does?

```
λ S . (if (singleton? S)
         "one"
         (next (L (select S)))).
```

One-knower

```
λ S . (if (singleton? S)
         "one"
         undef)
```

Two-knower

```
λ S . (if (singleton? S)
         "one"
         (if (doubleton? S)
             "two"
             undef))
```

Singular-Plural

```
λ S . (if (singleton? S)
         "one"
         "two")
```

Mod-5

```
λ S . (if (or (singleton? S)
              (equal-word? (L (set-difference S)
                              (select S)))
              "five"))
"one"
(next (L (set-difference S
                        (select S)))))
```

Three-knower

```
λ S . (if (singleton? S)
         "one"
         (if (doubleton? S)
             "two"
             (if (tripleton? S)
                 "three"
                 undef)))
```

CP-knower

```
λ S . (if (singleton? S)
         "one"
         (next (L (set-difference S
                              (select S)))))
```

2-not-1-knower

```
λ S . (if (doubleton? S)
         "two"
         undef)
```

2N-knower

```
λ S . (if (singleton? S)
         "one"
         (next (next (L (set-difference S
                              (select S)))))
```

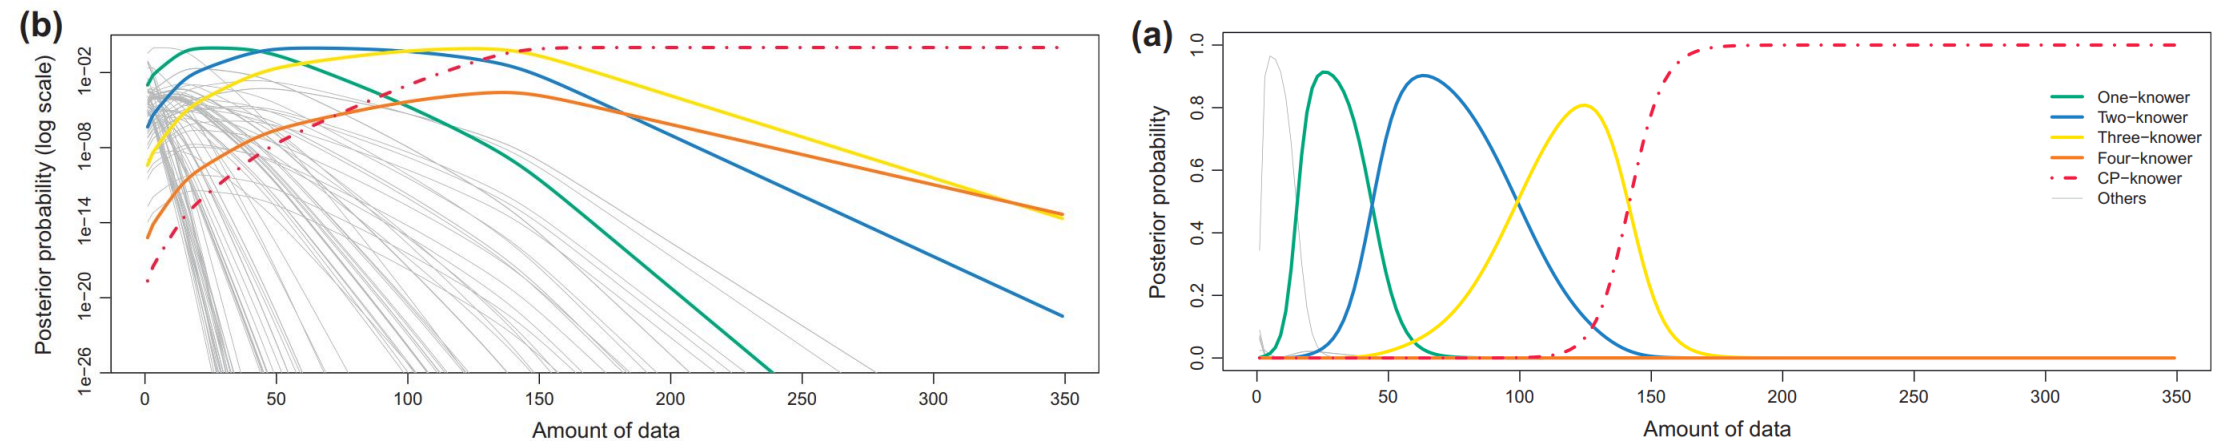
Learning Numerals with an LoT

- The likelihood function is pretty typical:
 - First, a set of objects is chosen from the universe of object, e.g. ‘cats’
 - Second, the hypothesis is evaluated on the set.
 - This can result either in a number word or ‘undef’
 - If the result is ‘undef’, a random number word is produced
 - If the result is a number word, the word is produced with probability α and with $1 - \alpha$ a random word is picked.
- Resulting likelihood function:

$$P(w_i|t_i, c_i, L) = \begin{cases} \frac{1}{N} & \text{if } L \text{ yields } \textit{undef} \\ \alpha + (1 - \alpha)\frac{1}{N} & \text{if } L \text{ yields } w_i \\ (1 - \alpha)\frac{1}{N} & \text{if } L \text{ does not yield } w_i \end{cases}$$

Learning Numerals with an LoT


- The model also penalizes recursive functions with a parameter γ
- Results look strikingly like human learning patterns:




- In particular, note all the hypotheses that are considered and then disregarded!

Learning Numerals with an LoT

- The same LoT can also learn things like singular/plural morphology:




- And mod-n systems (e.g. days of the week)




Visual Concept Learning in an LoT

- Overlan et al (2017), *Learning abstract visual concepts via probabilistic program induction in a Language of Thought*
- Last week, we have seen a model that can learn and do various other things with handwritten characters.
- Now let's see if we can work with something else in the visual modality, namely the structure of 3-d objects.
- Suppose we have the following primitive objects, and we can combine them in various ways:



Visual Concept Learning in an LoT


- For instance, we can produce the following categories:




- Note that each category identifies a *class* of objects!

Visual Concept Learning in an LoT

ABC



Ring



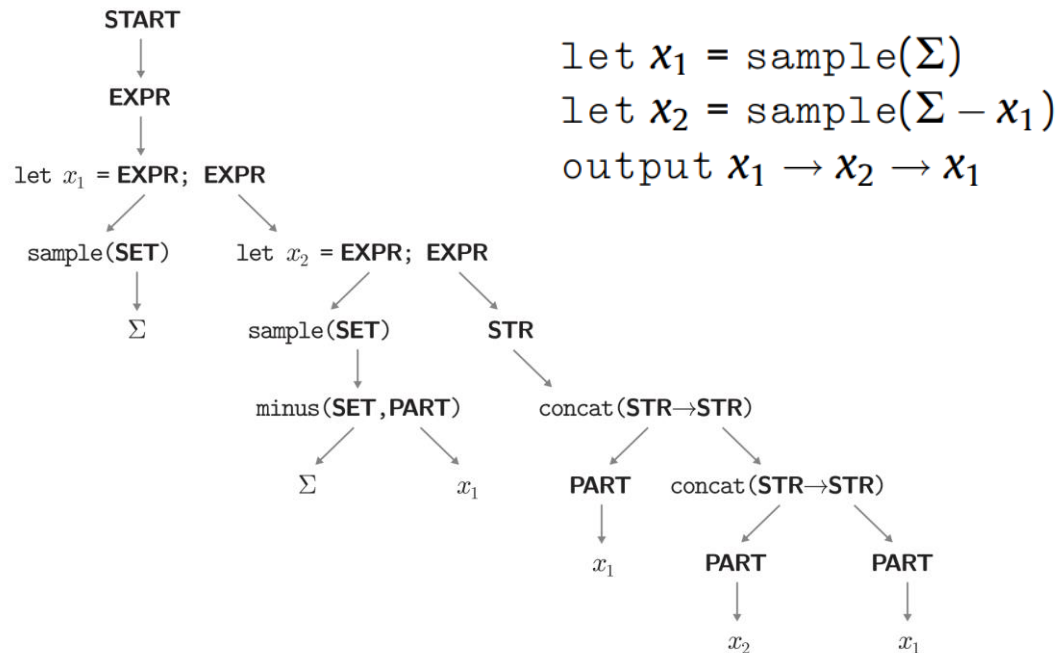
Visual Concept Learning in an LoT

- What could a grammar for this look like?

START	→ let <BV_PART>:x ₁ = FIRST_PART; EXPR	
EXPR	→ let <BV_PART>:x _n = PART; EXPR	
	→ STRING	
STRING	→ BV_PART	
	→ STRING CONNECT STRING	
	→ {STRING}	
FIRST_PART	→ sample(FIRST_SET)	1 - p _{single}
	→ SINGLE	p _{single}
PART	→ BV_PART	(1 - p _{single})/2
	→ sample(SET)	(1 - p _{single})/2
	→ SINGLE	p _{single}
FIRST_SET	→ Σ	1 - p _{minus}
	→ minus(FIRST_SET, FIRST_PART)	p _{minus}
SET	→ Σ	1 - p _{minus}
	→ minus(SET, BV_PART)	p _{minus}
CONNECT	→ '↑' '↓' '←' '→'	
SINGLE	→ 'a' ... 'e'	

Visual Concept Learning in an LoT

- Example of a derivation and some categories:



$\text{let } x_1 = \text{sample}(\Sigma)$
 $\text{let } x_2 = \text{sample}(\Sigma - x_1)$
 $\text{output } x_1 \rightarrow x_2 \rightarrow x_1$

ABA

$\text{let } x_1 = \text{sample}(\Sigma_R)$
 $\text{let } x_2 = \text{sample}(\Sigma_R - x_1)$
 $\text{output } x_1 \rightarrow x_2 \rightarrow x_1$

xBB

$\text{let } x_1 = \text{sample}(\Sigma)$
 $\text{let } x_2 = 'a'$
 $\text{output } x_2 \rightarrow x_1 \rightarrow x_1$

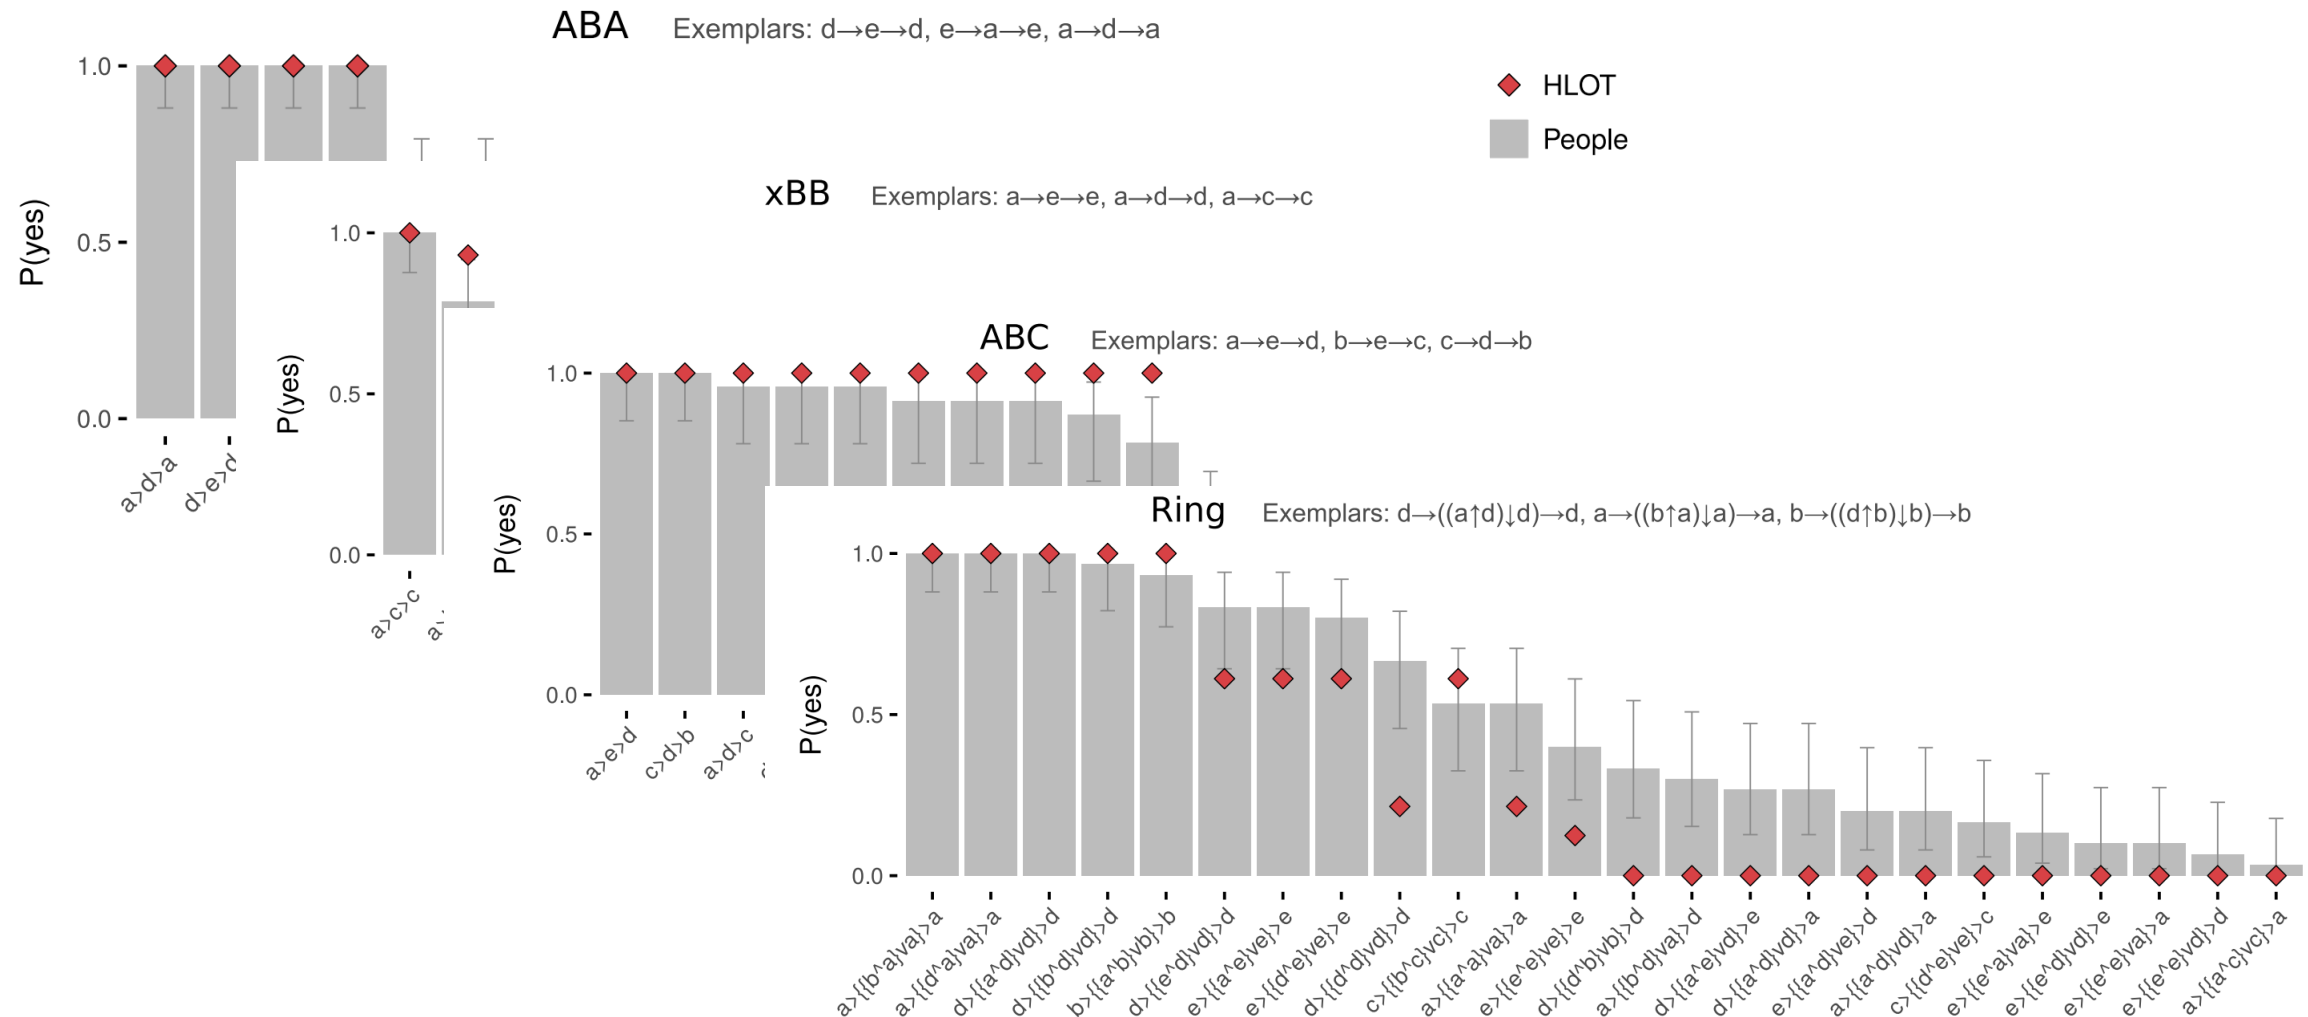
ABC

$\text{let } x_1 = \text{sample}(\Sigma)$
 $\text{let } x_2 = \text{sample}(\Sigma - x_1)$
 $\text{let } x_3 = \text{sample}(\Sigma - x_2 - x_1)$
 $\text{output } x_2 \rightarrow x_3 \rightarrow x_1$

Ring

$\text{let } x_1 = \text{sample}(\Sigma_R)$
 $\text{let } x_2 = \text{sample}(\Sigma_R - x_1)$
 $\text{output } x_1 \rightarrow ((x_2 \uparrow x_1) \downarrow x_1) \rightarrow x_1$

Visual Concept Learning in an LoT




Sequence Learning in an LoT

- Planton et al (2021), *A theory of memory for binary sequences: Evidence for a mental compression algorithm in humans*
- The scope here is to understand how humans deal with *binary sequences*, i.e. sequences composed of only two elements.
 - For instance, $\{0,1\}^*$
 - But not that it doesn't need to be symbols. E.g. it can be two pitches.
- Much literature has been devoted to understand how humans learn these.
- Usually, the general strategy is to find a mechanism where strings can be encoded, which explains how long strings can be learned, which would be impossible with pure memorization.

Sequence Learning in an LoT

- In this case, we'll use an LoT where each sentence can encode a binary series. For instance:



LoT program expression :


```
[[[+0]^2]^2<b>,[b]^4]^2<+0>
```

LoT complexity = 12

- LoT:
 - Staying (“+o”)
 - Moving to the other item (here denoted ‘b’)
 - Repetition (“^n”, where n is any number),
 - Possibly with a variation in the starting point
 - Denoted by <x> where x is an elementary instruction, either +o or b
 - Embedding of expressions is represented by brackets (“[. . .]”) and concatenation by commas (“,”).

Sequence Learning in an LoT

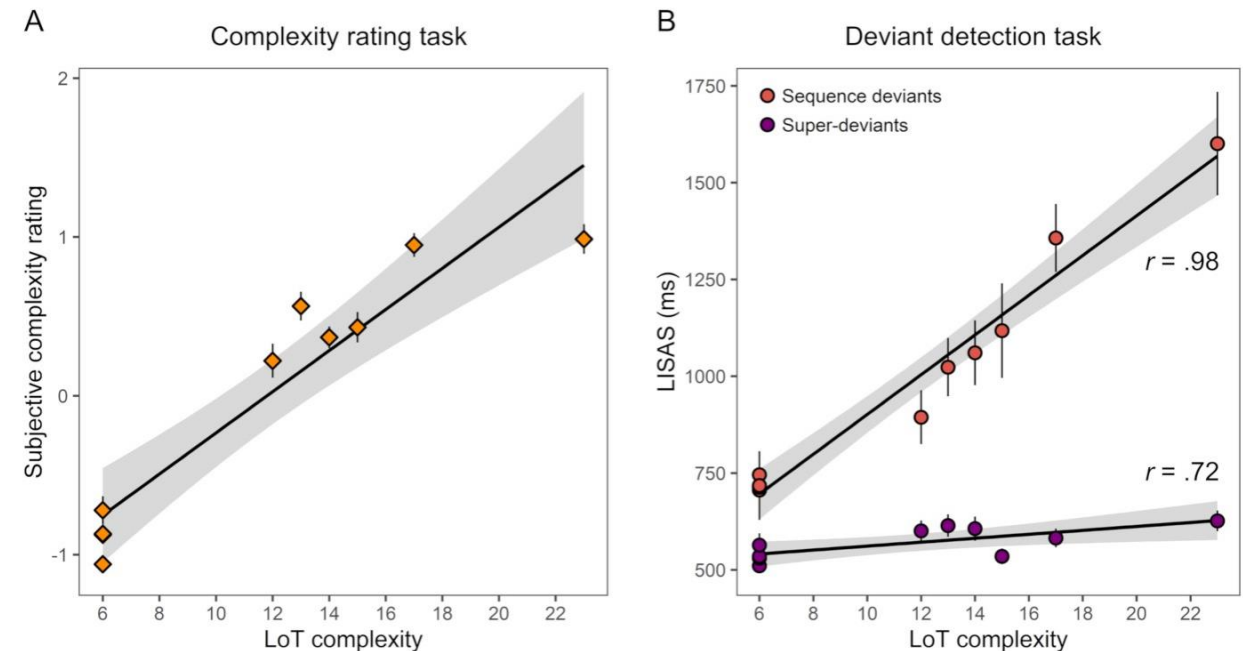
- ‘Sequence violation’ experimental paradigm:



- The basic hypothesis was that, for equal sequence length, error rate and response time in violation detection would increase with sequence complexity.

Sequence Learning in an LoT

- This is the results of the first experiment
- Participants were asked subjective complexity (left)
- And to identify deviations in:
- Sequence deviants
 - A note is replaces with the other note
- Superdeviant
 - A new note is introduced
- LISAS: Linear Integrated Speed-Accuracy Score



Summary

- This week, we saw some other applications of the pLoT idea
- First, we saw a model of numeral acquisition
- Then, a model of learning visual concepts (categories of 3-d objects)
 - This could easily be extended, if you're into 3-d rendering
 - Possible final project?
- Finally, we saw a model of learning binary sequences of tones
- In the lab this week, we'll try to implement a model from scratch
 - I am not going to prepare beforehand so we can all write it together
- Next week is the last week, so the lecture is going to be a review of what we've done, with some concluding remarks on the general project.
- In the lab next week we'll implement another model.